

TDT4136 Introduction to Artificial Intelligence

Lecture 1: Introduction (Chapter 1 in the textbook)

Pinar Öztürk

Norwegian University of Science and Technology

All practical Information is in the first part (separate) of Lecture 1 slides.

Course topic overview

- Intelligent agents
- Problem solving by Searching
- Adversarial Search
- Constraint Satisfaction Problems
- Logical systems
- Knowledge representation
- Planning
- Game Theory
- Ethical issues in AI

Today's outline

What is AI?

A brief history

The state of the art

What is AI?

There are no crisp definitions. Here is one from [John McCarthy](#), (Father of the phrase *Artificial Intelligence*)

see <http://www.formal.Stanford.EDU/jmc/whatisai/>

Question: What is artificial intelligence?

Answer: It is the science and engineering of making intelligent machines, especially intelligent computer programs. It is related to the similar task of using computers to understand human intelligence, but AI does not have to confine itself to methods that are biologically observable.

Question: Yes, but what is intelligence?

Answer: Intelligence is the computational part of the ability to achieve goals in the world. Varying kinds and degrees of intelligence occur in people, many animals and some machines.

What is AI? - cont

"It is the science and engineering of making intelligent machines, especially intelligent computer programs" [John McCarthy](#)

"AI is the science of making machines to do things that would require intelligence if done by men." [Marvin Minsky](#)

"The study of mental faculties through the use of computational models" [Eugen Charniak](#)

In short: There is no formal definition covering all aspects of intelligence

Two aspects of AI

Science

The science of understanding intelligent entities, developing theories which attempt to explain and predict the nature of such entities

Discover ideas about knowledge that help explain various sorts of intelligence

Model functions of the human brain

Engineering

Solving real-world problems by employing ideas of how to represent and use knowledge

Engineering of intelligent entities

Produce intelligent behaviour by any means


Can machines be Intelligent?

- *Symbolic system hypothesis* (Newell and Simon)
 - Intelligence is substrate neutral
 - A *physical symbol system* has necessary and sufficient means for general intelligent action.
- Biological substrate only (John Searle, philosopher)
 - Intelligence is substrate dependent.
 - The material humans are made of is fundamental for our intelligence.
 - Thinking is possible only in special machines - living ones made of proteins.


AI prehistory - grounding disciplines

Computer Science is the main discipline underlying AI

Can we think which other disciplines AI is grounded on?



AI prehistory - grounding disciplines



Philosophy	logic, methods of reasoning mind as physical system foundations of learning, language, rationality
Mathematics	logic, formal representation algorithms, computation, probability theory
Statistics	modeling uncertainty, learning from data
Psychology	phenomena of perception and motor control cognitive psychology

Economics	formal theory of rational decisions, utility
Linguistics	knowledge representation grammar, syntax, semantics
Neuroscience	neurons as information processing units synapse as learning mechanism
Control theory	homeostatic systems, stability simple optimal agent designs, maximize objective function
Computer science	engineering, hardware computational complexity theory

Birth of AI field

A Summer Research Project in the Dartmouth College, in 1956 was the birth of the AI research field.



(From left: Trenchard More, John McCarthy, Marvin Minsky, Oliver Selfridge, and Ray Solomonoff)

McCarthy: "...to proceed on the basis of the conjecture that every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it."

Brief history of AI

- 1943 McCulloch & Pitts: Boolean circuit model of brain
- 1950 Turing's "Computing Machinery and Intelligence"
- 1952-69 Early AI programs, including Samuel's checkers program, Newell & Simon's Logic Theorist, Gelernter's Geometry Engine
- 1956 Dartmouth meeting: "Artificial Intelligence" adopted
- 1962 Rosenblatt's Perceptron for training simple neural networks
- 1965 Robinson's complete algorithm for logical reasoning
- 1972 The logic programming language PROLOG is created.
- 1966-74 Disappointment: AI discovers computational complexity
Neural network research almost disappears

Brief history of AI -cont.

1969–79	Early development of knowledge-based systems
1980–88	Expert systems industry booms
1988–93	Expert systems industry busts: “AI Winter”
1985–95	Backpropagation learning returns neural networks to popularity
1988–	Resurgence of probability; general increase in technical depth “Nouvelle AI”: ALife, GA
1995–	Agents, agents, everywhere . . .
2003–	Human-level AI back on the agenda
2005-2010	2005-2010 AI disappoints again, AI is not much appreciated

Example from Early History of AI

1943 McCulloch & Pitts: Boolean circuit model of brain

BULLETIN OF
MATHEMATICAL BIOPHYSICS
VOLUME 5, 1943

A LOGICAL CALCULUS OF THE IDEAS IMMANENT IN NERVOUS ACTIVITY

WARREN S. MCCULLOCH AND WALTER PITTS

FROM THE UNIVERSITY OF ILLINOIS, COLLEGE OF MEDICINE,
DEPARTMENT OF PSYCHIATRY AT THE ILLINOIS NEUROPSYCHIATRIC INSTITUTE,
AND THE UNIVERSITY OF CHICAGO

Because of the "all-or-none" character of nervous activity, neural events and the relations among them can be treated by means of propositional logic. It is found that the behavior of every net can be described in these terms, with the addition of more complicated logical means for nets containing circles; and that for any logical expression satisfying certain conditions, one can find a net behaving in the fashion it describes. It is shown that many particular choices among possible neurophysiological assumptions are equivalent, in the sense that for every net behaving under one assumption, there exists another net which behaves under the other and gives the same results, although perhaps not in the same time. Various applications of the calculus are discussed.

1. Introduction

Theoretical neurophysiology rests on certain cardinal assumptions. The nervous system is a net of neurons, each having a soma and an axon. Their adjunctions, or synapses, are always between the axon of one neuron and the soma of another. At any instant a neuron has some threshold, which excitation must exceed to initiate an impulse. This, except for the fact and the time of its occurrence, is determined by the neuron, not by the excitation. From the point of excitation the impulse is propagated to all parts of the neuron. The velocity along the axon varies directly with its diameter, from less than one meter per second in thin axons, which are usually short, to more than 150 meters per second in thick axons, which are usually long. The time for axonal conduction is consequently of little importance in determining the time of arrival of impulses at points unequally remote from the same source. Excitation across synapses oc-

- Goal:** to understand how the brain produces complex thoughts ("propositions") by using (simple) neurons
- model of a neuron, axon, dendrode - "MCP neuron"
 - network of neurons
 - transfer of information through on/off mechanism of neurons
 - every network of MCP neurons encodes some logical proposition

McCulloch and Pitts 1943 - "Bird example"

Example from:


http://www.mind.ilstu.edu/curriculum/mcp_neurons/mcp_neuron_1.php?modGUI=212&compGUI=1749&itemGUI=3018

Assume a bird will "decide" to eat/not an object (e.g., blue berry, orange, basketball, daisy). It receives/perceives two inputs about the object:

- shape: round or not
- colour: purple or not
- decision: eat only if round purple object (e.g., blue berry)

Object	Purple?	Round?	Eat?
Blueberry	Yes	Yes	Yes
Golf ball	No	Yes	No
Violet	Yes	No	No
Hot Dog	No	No	No

Bird example - cont.



- two inputs (each takes value 0 or 1)
- threshold T
- output is a function of inputs and T
- **IF**(sum(inputs)) $\geq T$
THEN Output=1

Object	Purple?	Round?	Eat?
Blueberry	1	1	1
Golf ball	0	1	0
Violet	1	0	0
Hot Dog	0	0	0

Object	Purple?	Round?	Total	Greater than or equal to threshold of 1?	Eat?
Blueberry	1	1	2	Yes	1
Golf ball	0	1	1	Yes	1
Violet	1	0	1	Yes	1
Hot Dog	0	0	0	No	0

Some killer apps in AI

1991 During the Gulf War, US forces deployed an AI logistics planning and scheduling program that involved up to 50,000 vehicles, cargo, and people. Saved the US more money than spent on all AI research since 1950

1997 Deep Blue (IBM) defeated world chess champion Gerry Kasparov

2011 Watson (IBM) beat human champions on “Jeopardy”


2012 Google car obtains driver's license in Nevada, US. By 2014, the cars have driven for 1.1 million km without accident

2017 DeepMind's Alphago AND Elon Musk's A.I. Destroys Champion Gamer! <https://www.youtube.com/watch?v=XbDmxEOj90Y>

What is AI?

Russell and Norvig's definition of AI

Two dimensions:



What is AI?

Types/approaches to AI - as defined by Russel and Norvig:

	humanly	rational
think		
act		

How R&N define AI

humanly vs. rationally

thinking
vs.
acting

Systems that think like humans	Systems that think rationally
Systems that act like humans	Systems that act rationally Rational Agents

1960s “**cognitive revolution**”: information-processing psychology replaced prevailing orthodoxy of **behaviorism**.

Tries to form computational theories of internal activities of the brain.

How to validate? Requires

- 1 Predicting and testing behavior of human subjects (top-down), or
- 2 Direct identification from neurological data (bottom-up)

Both approaches (roughly, **Cognitive Science** and **Cognitive Neuroscience**) are now distinct from AI

Both share with AI the following characteristic:
the available theories do not explain (or engender)
anything resembling human-level general intelligence

John McCarthy talks about AI


<https://www.youtube.com/watch?v=0zipf13jRr4>

Acting humanly: The Turing test

Turing (1950) “Computing machinery and intelligence”:

Can machines think?” → “**Can machines behave intelligently?**”

Operational test for intelligent behavior: the **Imitation Game**



- One of the most famous early AI programs: Eliza, the computer psychotherapist, created by Joseph Weizenbaum in 1966 at MIT.
- Eliza functions by "twisting the statements of her 'patients' back at them in the classic manner of a non-directive psychotherapist."
- The fact that it understands the subject's statements is an illusion.
- Surprisingly, many users were taking its performance quite seriously.

Acting Humanly - Eliza 2

young woman: Men are all alike.

eliza: In what way?

young woman: They're always bugging us about something specific or other.

eliza: Can you think of a specific example?

young woman: Well, my boyfriend made me come here.

eliza: Your boyfriend made you come here?

young woman: He says I'm depressed much of the time.

eliza: I am sorry to hear you are depressed.

young woman: It's true. I'm unhappy.

eliza: Do you think coming here will help you not to be unhappy?

young woman: I need some help, that much seems certain.

eliza: What would it mean to you if you got some help?

young woman: Perhaps I could learn to get along with my mother.

eliza: Tell me more about your family.

young woman: My mother takes care of me.

eliza: Who else in your family takes care of you?

young woman: My father.

eliza: Your father.

young woman: You are like my father in some ways.

eliza: What resemblance do you see?

young woman: You are not very aggressive but I think you don't want me to notice that.

eliza: What makes you think I am not very aggressive?

young woman: You don't argue with me.

eliza: Why do you think I don't argue with you?

young woman: You are afraid of me.

eliza: Does it please you to believe I am afraid of you?

young woman: My father is afraid of everybody.

eliza: What else comes to mind when you think of your father?

young woman: Bullies.

eliza: Does that have anything to do with the fact that your

> Hello. My name is Eliza. How may I help you?

Eliza Chat bot

Is she a Rogerian psychotherapist, a semi intelligent chat bot, or just a toy?



Eliza, one of the original chatbots, is now ready to talk to you.

She is programmed to talk like a Rogerian psychotherapist, an interesting example of the limitations of early artificial intelligence programs.

If Eliza (or you, or your web browser) gets stuck, [refresh the page](#). To start a new session with the therapist [reload the page](#).

How humany is Sofia (Hong Kong Kong firm Hanson Robotics)?

<https://www.youtube.com/watch?v=suRuQbDXcrc>

Thinking rationally: Laws of Thought-1

- Aristotle: what are correct arguments/thought processes?
Formalize “correct” reasoning using a mathematical model
- Several Greek schools developed various forms of **logic**:
notation and *rules of derivation* for thoughts;
- Direct line through mathematics and philosophy to modern AI

Thinking rationally: Laws of Thought-2



Thinking rationally: Laws of Thought-3



Problems:

- 1) Not all intelligent behavior is mediated by logical deliberation
- 2) What is the purpose of thinking? What thoughts *should* I have out of all the thoughts (logical or otherwise) that I *could* have?
- 3) Formalizing (informal) common sense knowledge is difficult
- 4) General deductive inference is computationally intractable

Acting rationally: Rational Agents approach

- **Rational** behavior: doing the right thing
- The right thing: that which is expected to maximize goal achievement, given the available information and computational abilities
- Doesn't necessarily involve thinking—e.g., blinking reflex—but thinking should be in the service of rational action

Two main AI paradigms

Good old fashioned AI (GOFAI)

Situated embodied AI (SEAI)

- Most of the successes of AI in the 1970s and 1980s were due to research based on Newell's and Simon's *Physical Symbol System Hypothesis*
- A physical symbol system has the necessary and sufficient means for general intelligent action.
- Newell and Simon viewed intelligence as *symbol manipulation*, and hypothesized that it didn't make difference what physical medium – brain, paper, or computer – was used to do the symbol manipulation
- Hence a special emphasis on *symbolic representations*, which can be interpreted as representing situations in the real world, e.g. "a block world"

- **Searle:** a program (or any physical symbol system) could not be said to understand the symbols that it uses; the symbols have no meaning for the machine
- **Brooks:** our most basic skills of motion, survival, perception, balance etc. do not seem to require high level symbols at all; the use of high level symbols was more complicated and less successful
- **Harnad:** the **symbol grounding** problem: an agent does not perceive symbols, instead the brain converts sensory inputs into higher level abstractions, e.g. symbols

- The GOFAI approaches turned out to be brittle and very little robust when deployed on real-world problems
- Trying to define a model of the world turned out to be quite hard - this led to Brooks' statement that "the world is its own best model"
- Situated and embodied AI focuses on having a body (i.e. motor skills) in a physical environment
- Swarm intelligence, subsymbolic AI, genetic algorithms, neural networks - however, this is not a big part of this course, this is covered in more detail in other AI courses.

Next Lecture: Intelligent Agents (chapter 2 in R&N book)

Computers with common sense(Doug Lenat):

https://www.youtube.com/watch?v=2w_ekB08ohU

Artificial Intelligence: The Common Sense Problem (Hubert Dreyfus)

<https://www.youtube.com/watch?v=SUZUbYCBtGI>

AI Boom (Rodney Brooks):

<https://www.youtube.com/watch?v=ovbzjLobiDo>

IBM Watson : https://www.youtube.com/watch?v=_Xcmh1LQB9I

Jeopardy : https://www.youtube.com/watch?v=WFR3l0m_xhE

State of the art AI (Minsky):

https://www.youtube.com/watch?v=a0DnFdU_hds