

SSDs, LSM-trees and RocksDB

Håvard Dybvik

Svein Erik Bratsberg, IDI/NTNU

Contents



- HDDs
- SSDs
- Sequential writes
- B+-trees
- LSM-trees
- RocksDB
- MyRocks
- Write amplification
- Write stalls and write stops

HDDs



- Rotating, magnetic disks
- Have been developed since 1956
- Storing vast amounts of data at low costs
- Access time does not improve much with new disks
- 5-10 millisecs today, the first (1956) had 600 millisecs.
- Throughput (typical every day, desktop disk):
 - 160 MB/s (write)
 - 180 MB/s (read)
- Needs to be careful with layout of file system

SSDs (1)




- Solid state drives – purely electronic devices
- Faster access times, lower latency, lower power consumption, completely silent operation, uniform random access speed
- NAND: Block: 64 or 128 *pages* of 2 Kbyte or 4 Kbyte
Block: 128 KB to 512 KB
- Reads and writes of *pages*
- Erase *complete blocks*
- Erase-before-write
- Throughput: Four times of HDD (at least), 550 MB/s

SSDs – wear leveling

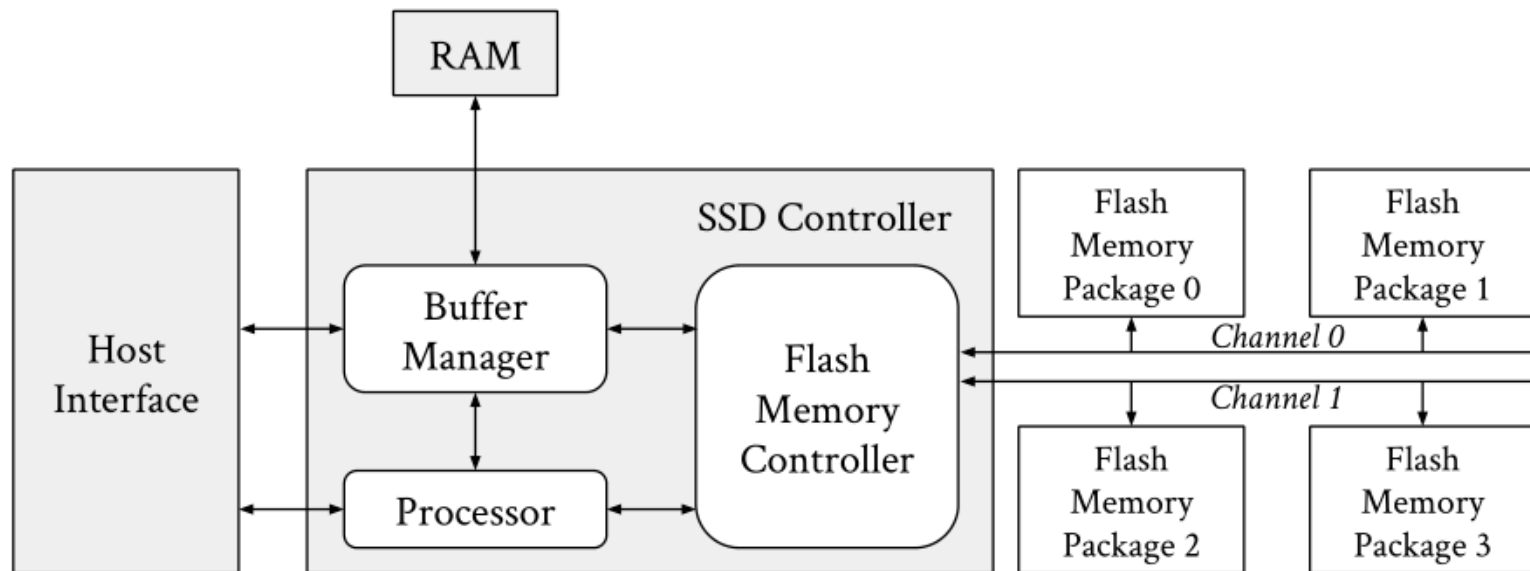


- Limited number of writes possible (wearing)
- Flash Translation Layer – firmware layer implementing wear leveling: Standardized by Intel.
 - *No wear leveling*: Fixed mapping from logical addresses to physical addresses
 - *Dynamic wear leveling*: At updates, the old block is marked as invalid and the block is relocated.
 - *Static wear leveling*: Also moves static blocks periodically. All parts of the disk will be worn out eventually.
- Garbage collection is important in SSDs. Done in units of blocks.
- Host based FTL (computer) and array-based FTL (disk)

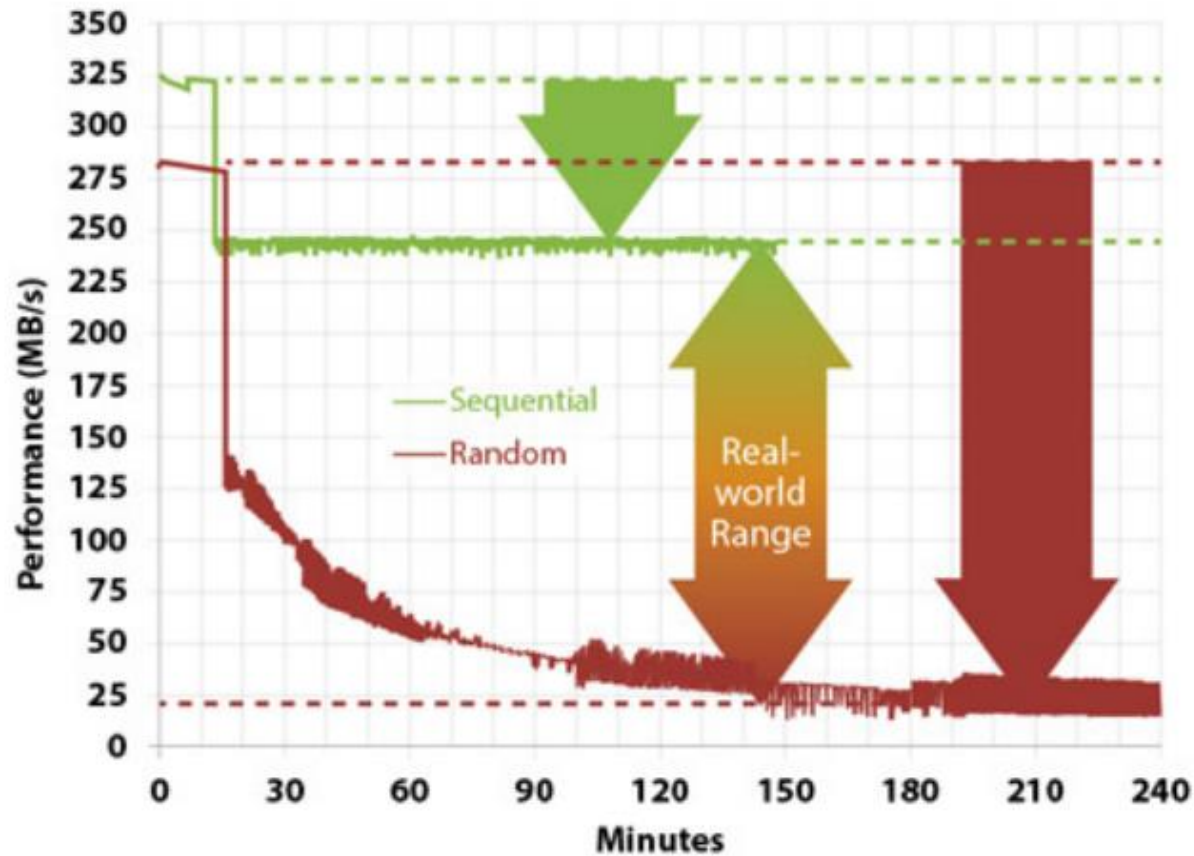
SSDs – garbage collection





SSDs – parallel I/O



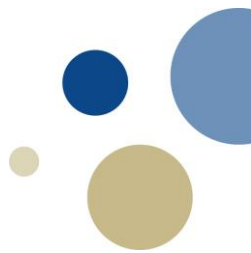
Sequential writes (1)



Sequential writes (2)



Sequential writes (3)



- One of the essential keys to great write performance is to organize data in such a way that each write request distributes the overhead of the write over multiple inserts batched together – large blocks are also good
- *Over-provisioning*: When the GC has too much to do, writes happen more often than erase operations -- The SSD has a reserved area where the garbage collector could put writes at high peaks periods. 7 – 28 % of disk.

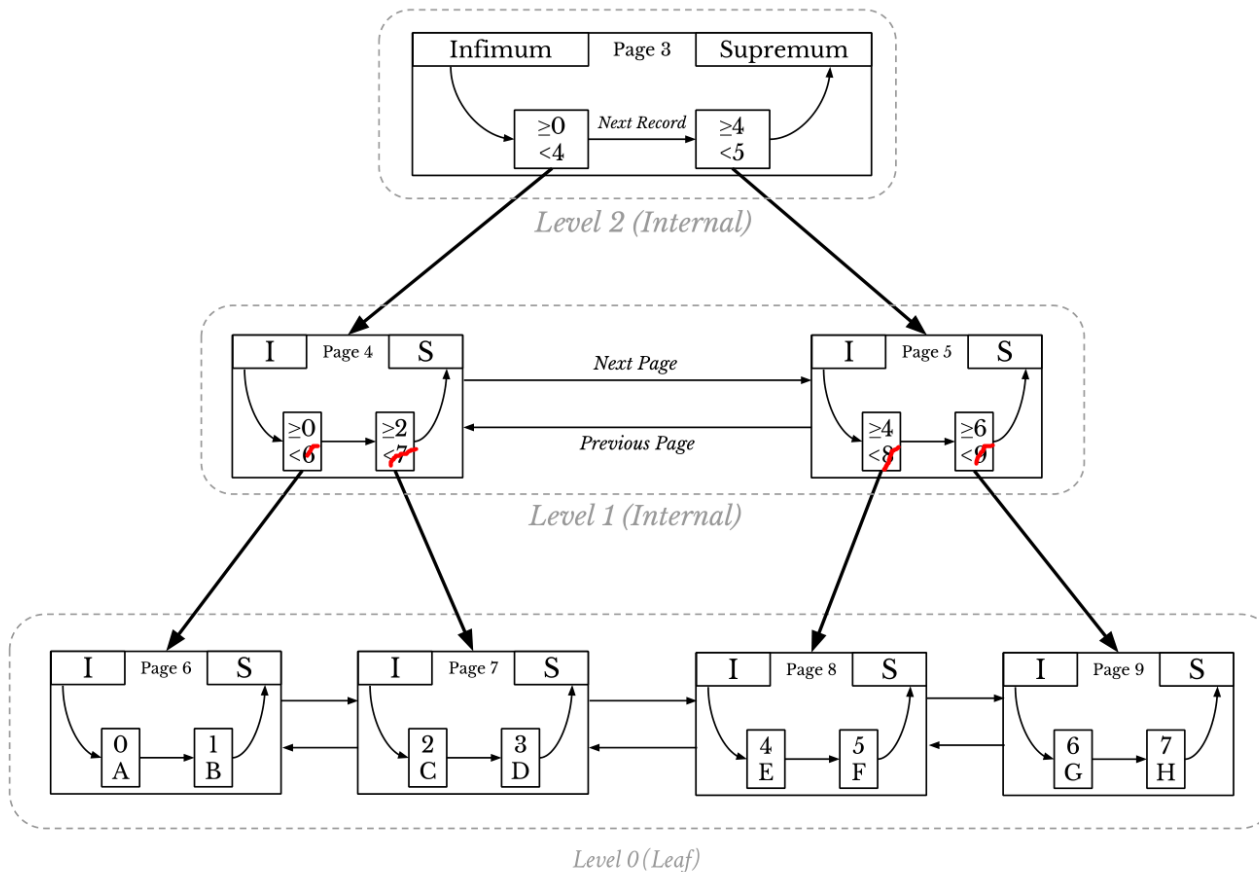
Writes/reads measured (old exercise)

| | HDD/w | HDD/r | SSD/w | SSD/r |
|------|-----------|-----------|----------|----------|
| 1GB | 1976 MB/s | 3835 MB/s | | |
| 2GB | 2101 MB/s | 3954 MB/s | | |
| 4GB | 1766 MB/s | 202 MB/s | | |
| 8GB | 386 MB/s | 157 MB/s | | |
| 16GB | 221 MB/s | 162 MB/s | | |
| 32GB | 184 MB/s | 163 MB/s | 354 MB/s | 463 MB/s |

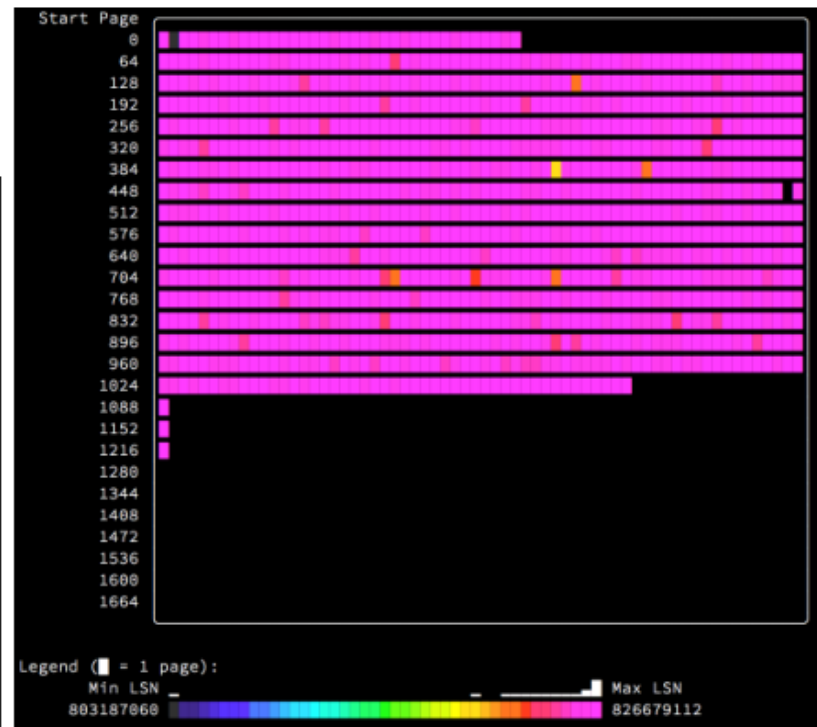

B+-trees (1)



B⁺-Tree Structure of InnoDB

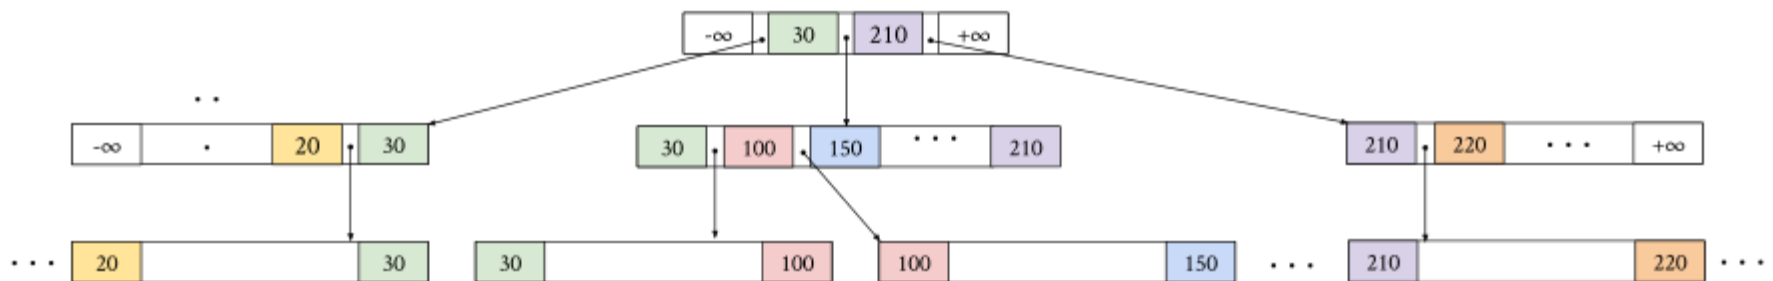


B+-trees (2) Freshness of PageLSNs





Variants of B+-trees

- *FB-tree*: Variable sized blocks simulating FTL's GC algorithm.
- *Copy-on-write B-tree*: Copy a block to a new location at updates. E.g. LMDB. BTRFS (file system on Linux).
- *Write-optimized B-tree*: No sideways pointers, index records contain lowKey and highKey of the leaf blocks. Prevents multiple writes due to pointers when moving a block.




LSM-trees



RocksDB

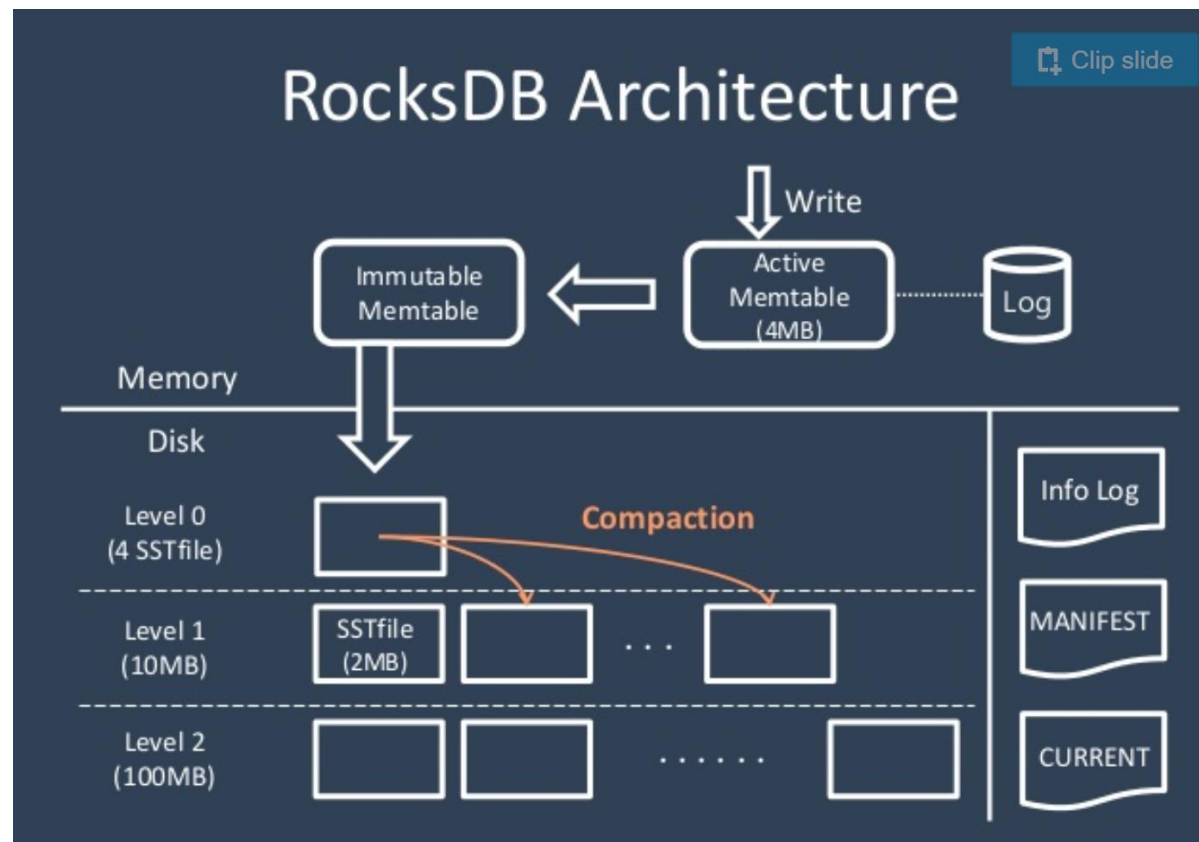


- Builds on Google's LevelDB
- Multithreaded compaction
- Multithreaded insertion into MemTable
- Extensive control over Bloom filters
- SSD support
- 10 x improved write performance compared to LevelDB




RocksDB (2)

- SSTable (sorted key ranges)
- Memtable
- Write Ahead Log




Leveled Compaction

- Original compaction style in LevelDB
- SSTable files stored in multiple levels
- Multiple overlapping SSTables at level 0 to get fast write of MemTables



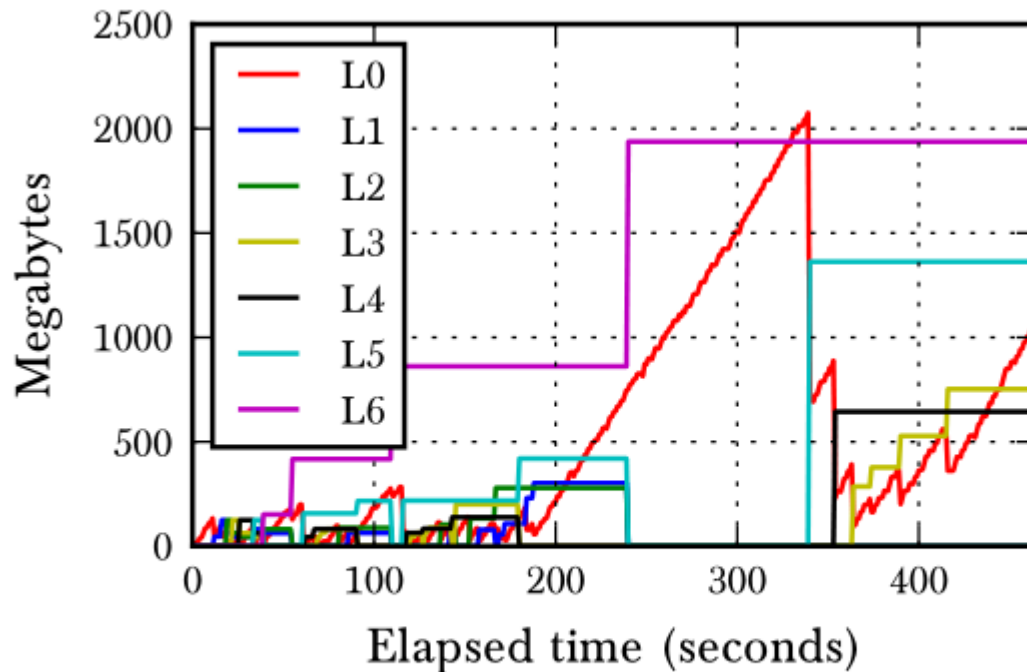
Leveled Compaction (2)



(a) Leveled compaction.

Universal Compaction

- SSTables overlap in key range, but not in time
- Merging SSTables to cover bigger time ranges



(b) Universal compaction.

Bloom filter

- Simple data structure to quickly check if a key *may* exist in a dataset
- Bitmap of m bits and k unique hash functions
- Either used on each block or the complete SSTables


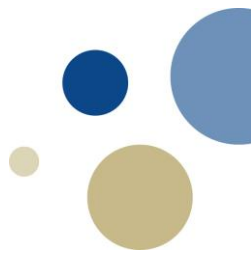


Figure 2.25: Example of a Bloom filter with the dataset x, y, z and $m = 15$ bits. There are $k = 3$ different hash functions.

MyRocks



- Facebook's integration of RocksDB as a storage engine with MySQL
- Created due to serious space and write amplification of MySQL's InnoDB
- MariaDB (MySQL competitor) is also integrating RocksDB as a storage engine

Write amplification



- $\text{\#bytes written to api} / \text{\#bytes written to database}$
- B-trees: write amp less with big records
- LSM-trees: less write amp due to big write chunks, but compaction require more writes
- LSM-trees: write amp independent on record size
- May be measured by insert throughput. RocksDB performs very well

Write stalls and write stops



- Write stalls: Slow the speed of inserts to cope with compaction
- Write stops: Stop inserts.
- RocksDB can configure parameters to control this:
 - Number of MemTables
 - Number of SSTables at level 0
 - Number of bytes awaiting compactions
- Hans-Wilhelm Kirsch Warlo developed an auto-tuner for compactions in RocksDB during his master thesis: Turns off compactions during high insert loads (spring 2018).

Warlo's auto-tuner

