



# Mining Streaming Data



# Goal

To understand the **importance** of mining methods for streaming data, and learn about existing **methods handling** streaming data



# Characteristics/Description

- Stream data sets are...
  - Continuous
  - Massive
  - Unbounded
  - Possibly infinite
- Fast changing and requires fast, real-time response
  - **Examples:**
    - Life threatening: collision avoidance
    - Lost revenue/transactions: hung-up networks



# Sizing the challenge



- WalMart Records >20 Million Transactions
- Google Handles >100 Million Searches
- AT&T produces >275 million call records
- Earth sensing satellite produces GBs of data
- Twitter handles millions of tweets in a second

This just in a day!

# The Stream Model

- Input tuples enter at a rapid rate, at one or more input ports.
- Impractical to store the whole streaming data
- How do you make critical calculations about the stream using a limited amount of (primary or secondary) memory?
- Simple calculation per data due to time and space constraints

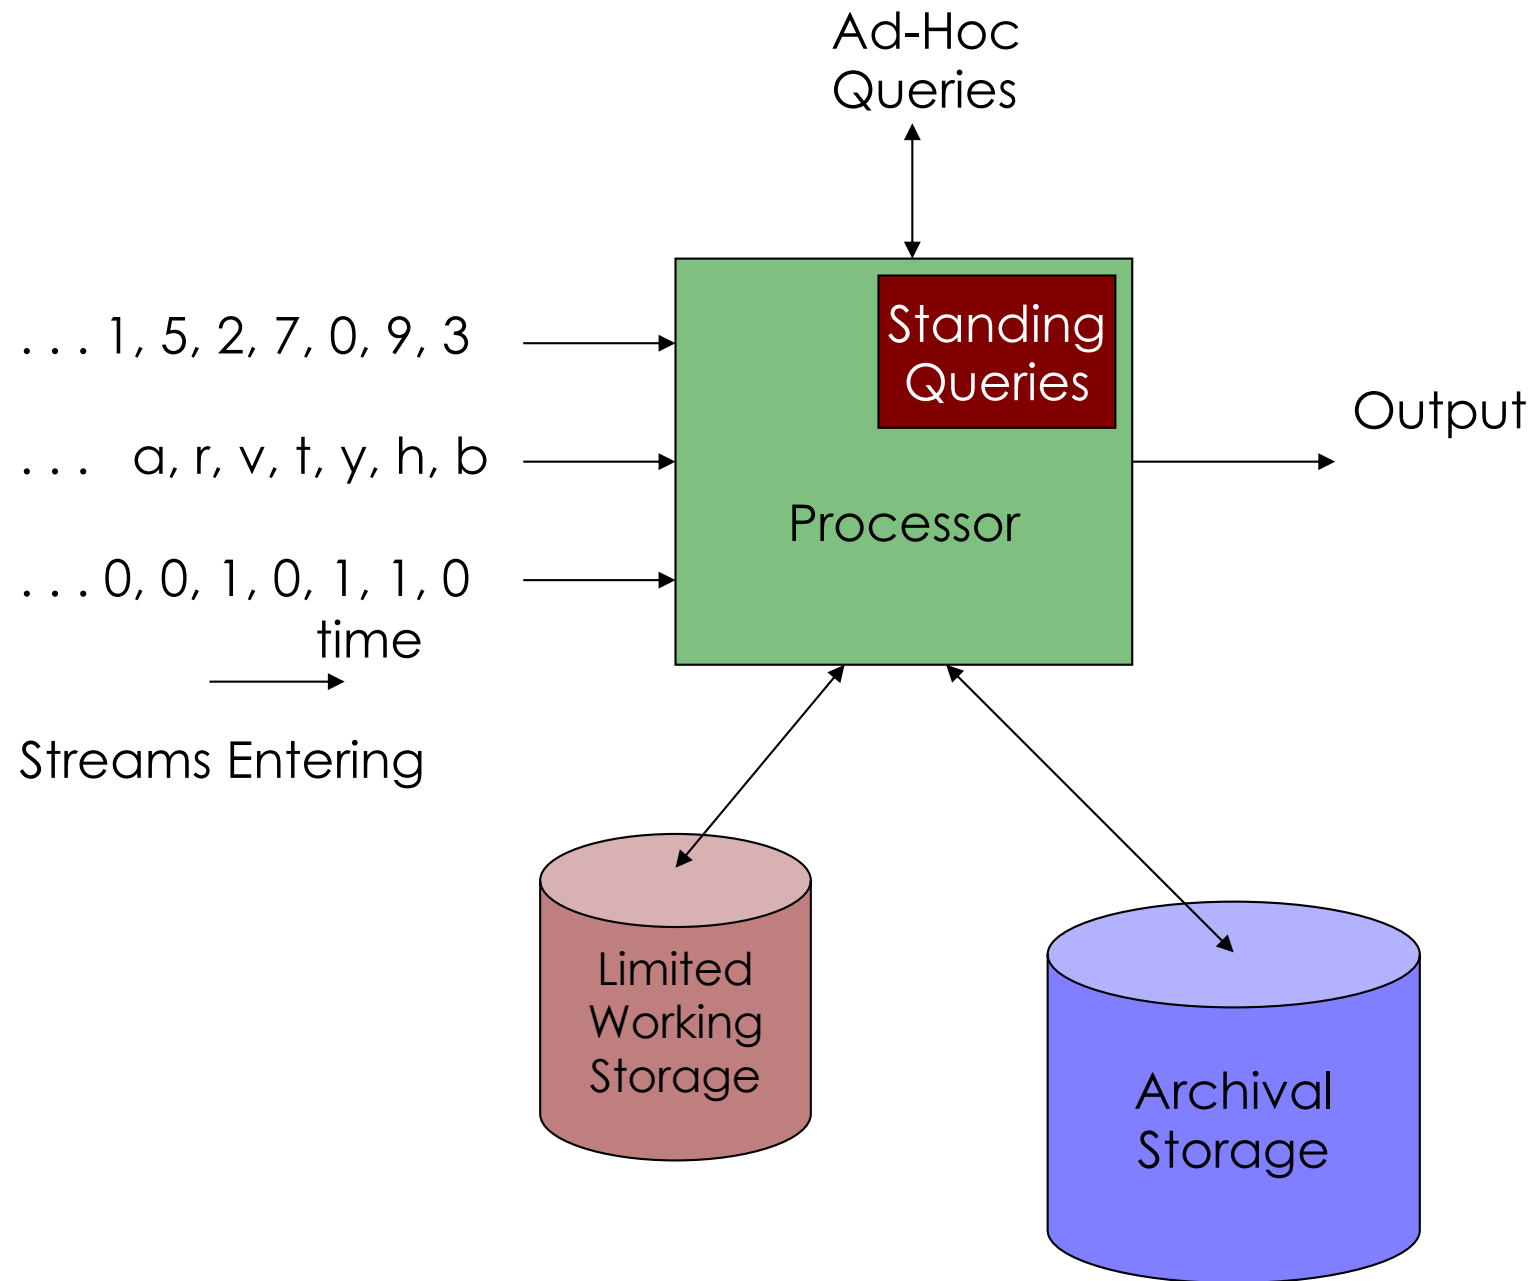
# Time/Space constrained

- Not enough memory
- Can't afford storing/revisiting the data
  - Single pass computation
- External memory algorithms for handling data sets larger than main memory cannot be used.
  - Do not support continuous queries
  - Too slow real-time response



# Two Forms of Query

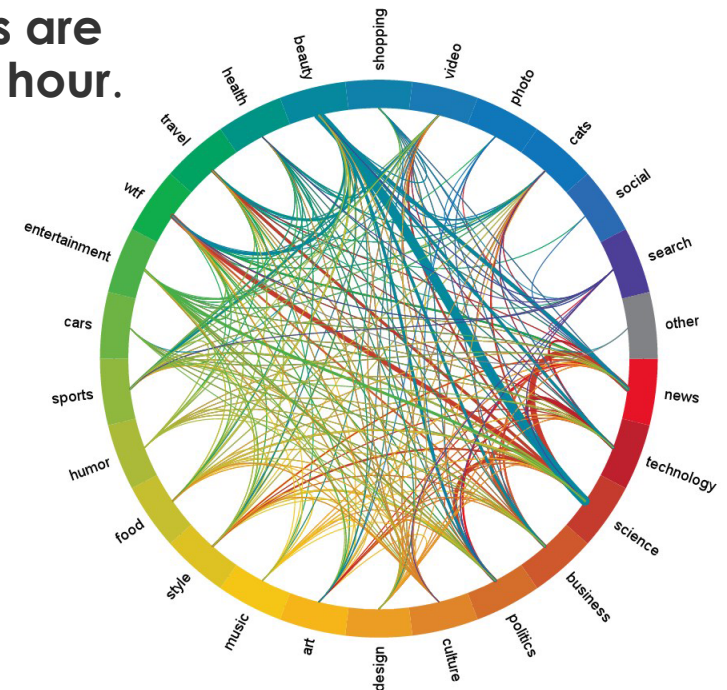
1. **Ad-hoc queries:** Normal queries asked one time about streams.
  - **Example:** What is the maximum value seen so far in stream  $S$ ?
2. **Standing queries:** Queries that are, in principle, asked about the stream at all times.
  - **Example:** Report each *new* maximum value ever seen in stream  $S$ .





# Example Applications

- Mining query streams.
  - E.g.: Google wants to know what queries are **more frequent today than yesterday**.
- Mining click streams.
  - E.g.: Yahoo wants to know **which of its pages are getting an unusual number of hits in the past hour**.
    - Often caused by annoyed users clicking on a broken page.
- IP packets can be monitored at a switch.
  - E.g.:
    - Gather information for optimal routing.
    - Detect denial-of-service attacks.



# Sliding Windows

- A useful model of stream processing
  - queries are about a **window** of length  $N$  – the  $N$  most recent elements received.
  - **Alternative**: elements received within a time interval  $T$ .
- **Interesting case**:
  - **$N$  too large** to be stored in main memory
  - Or, **too many streams** that windows for all do not fit in main memory.

# Sliding Windows

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

← Past Future →

## Example: Averages

- Stream of integers, window of size  $N$ .
- **Standing query:** what is the average of the integers in the window?
- For the first  $N$  inputs, sum and count to get the average.

### Question:

- What do we do now for the next input  $i$ ?

## Example: Averages

### Answer:

- When a new input  $i$  arrives:
  - change the average by adding  $(i - j)/N$ , where  $j$  is the oldest integer in the window before  $i$  arrived.
- **Good**:  $O(1)$  time per input.
- **Bad**: Requires the entire window in main memory.



# Counting 1's

Approximating Counts  
Exponentially Growing Blocks  
DGIM Algorithm

# Approximate Counting

- Can show that an exact sum or count of the elements in a window:
  - Cannot use less space than the window itself.
- If willing to accept an **approximation**:
  - Can use much less space.
- Consider the simple case of counting bits
- Sums are a fairly straightforward extension.

# Counting Bits

- **Problem:** given a stream of 0's and 1's, be prepared to answer queries of the form “how many 1's in the most recent  $k$  bits?” where  $k \leq N$ .
- **Obvious solution:** store the **most recent  $N$  bits**
- But answering the query will take  $O(k)$  time
  - Very possibly too much time.
- And the space requirements can be too great.
  - Especially if many streams to be managed in main memory at once, or  $N$  is huge.



## Example: Bit Counting

- Count recent hits on URL's belonging to a site.
- Stream is a sequence of URL's.
- Window size  $N = 1$  billion.
- Think of the data as many streams – one for each URL.
- Bit on the stream for URL  $x$  is 0 unless the actual stream has  $x$ .

# DGIM Method

- Name refers to the inventors:
  - **D**atar, **G**ionis, **I**ndyk, and **M**otwani.
- Store only  $O(\log_2 N)$  bits per stream.
  - $N$  = window size.
- Gives approximate answer, never off by more than 50%.
  - Error factor can be reduced to any  $\varepsilon > 0$ , with more complicated algorithm and proportionally more stored bits.

# Timestamps

- Each bit in the stream has a *timestamp*, starting 0, 1, ...
- Record timestamps modulo  $N$  (the window size), so we can represent any *relevant* timestamp in  $O(\log_2 N)$  bits.


# Buckets

- A segment of the window, represented by a record consisting of:
  - The **timestamp** of its end [ $O(\log N)$  bits].
  - The **number of 1's between its beginning and end**.
    - Number of 1's = **size** of the bucket.
- **Constraint on bucket sizes:** number of **1's must be a power of 2**.
  - Thus, only  $O(\log \log N)$  bits are required for this count.

# Representing a Stream by Buckets

- Either one or two buckets with the same power-of-2 number of 1's.
- Buckets do not overlap.
- Buckets sorted by size.
  - Older buckets not smaller than newer buckets.
- Buckets disappear when their end-time is  $> N$  time units in the past.

# Example: Bucketized Stream



# Updating Buckets

- New bit in, drop the last (oldest) bucket if its end-time is prior to  $N$  time units before the current time.
- If the current bit is 0, no other changes are needed.

# Updating Buckets: Input = 1

- If the current bit is 1:
  - Create a new bucket of size 1, for just this bit.
    - End timestamp = current time.
  - If there are now three buckets of size 1, combine the oldest two into a bucket of size 2.
  - If there are now three buckets of size 2, combine the oldest two into a bucket of size 4.
  - And so on ...



# Example: Managing Buckets

Initial

10010101100010110101010101010110101010101011101010101110101010111010101000101100010

1 arrives; makes third block of size 1.

00101011000101101010101010101101010101010111010101011101010101110101010001011000101

Combine oldest two 1's into a 2.

00101011000101101010101010101101010101010111010101011101010101110101010001011000101

Later, 1, 0, 1 arrive. Now we have 3 1's again.

0101100010110101010101010101101010101011101010101110101010001011000101101

Combine two 1's into a 2.

0101100010110101010101010101101010101011101010101110101010001011000101101

The effect ripples all the way to a 16.

0101100010110101010101010101101010101011101010101110101010001011000101101

# Querying

- To estimate the number of 1's in the most recent  $k < N$  bits:
  - Restrict your attention to only those buckets whose end time stamp is at most  $k$  bits in the past.
  - Sum the sizes of all these buckets but the oldest.
  - Add half the size of the oldest bucket.
- Remember:
  - we don't know how many 1's of the last bucket are still within the window.

# Error Bound

- Suppose the oldest bucket within range has size  $2^i$ .
- Then by assuming  $2^i - 1$  of its 1's are still within the window, we make an error of at most  $2^i - 1$ .
- Since there is at least one bucket of each of the sizes less than  $2^i$ , and at least 1 from the oldest bucket, the true sum is no less than  $2^i$ .
- Thus, error at most 50%.

# Space Requirements

- Can represent one bucket in  $O(\log N)$  bits.
  - It's just a timestamp needing  $\log N$  bits and a size, needing  $\log \log N$  bits.
- No bucket can be of size greater than  $N$ .
- There are at most two buckets of each size  $1, 2, 4, 8, \dots$
- That's at most  $\log N$  different sizes, and at most 2 of each size, so at most  $2\log N$  buckets.



- Bloom Filters
- Sampling Streams
- Counting Distinct Items

# Filtering Stream Content

- To motivate the **Bloom-filter** idea, consider a web crawler.
- It keeps, centrally, a list of all the URL's it has found so far.
- It assigns these URL's to any of a number of parallel tasks
  - these tasks stream back the URL's they find in the links they discover on a page.
- It needs to filter out those URL's it has seen before.

# Role of the Bloom Filter

- A Bloom filter placed on the stream of URL's will declare that **certain URL's have been seen before**.
- Others will be declared new, and will be added to the list of URL's that need to be crawled.
- Unfortunately, the Bloom filter can have false positives.
  - It can declare a **URL seen before when it hasn't**.
- But if it says "never seen," then it is truly new.
- **Solution:** restart the filter periodically.

## Example: Filtering Chunks

- Suppose we have a database relation stored in a DFS, spread over many chunks.
- We want to find a particular value  $v$ , looking at as few chunks as possible.
- A Bloom filter on each chunk will tell us certain values are there, and others aren't.
  - As before, **false positives possible**
- But now things are exactly right: if the filter says  $v$  is not at the chunk, it surely isn't.
  - Occasionally, we retrieve a chunk we don't need, but can't miss an occurrence of value  $v$ .



# How a Bloom Filter Works

- A *Bloom filter* is an array of bits, together with a number of hash functions.
- The argument of each hash function is a stream element, and it returns a position in the array.
- Initially, all bits are 0.
- When input  $x$  arrives, we **set to 1 the bits  $h(x)$ , for each hash function  $h$ .**

## Example: Bloom Filter

- Use  $N = 11$  bits for our filter.
- Stream elements = integers.
- Use two hash functions:
  - $h1(x) =$ 
    - Take **odd-numbered bits from the right** in the binary representation of  $x$ .
    - Treat it as an integer  $i$ .
    - Result is  $i$  modulo 11.
  - $h2(x) =$  same, but **take even-numbered bits**.

## Example – Continued

Stream element	$h_1$	$h_2$	Filter contents
			000000000000
25 = 11001	5	2	001001000000
159 = 10011111	7	0	101001010000
585 = 1001001001	9	7	10100101010

Note: bit 7 was already 1.

# Bloom Filter Lookup

- Suppose element  $y$  appears in the stream, and we want to know if we have seen  $y$  before.
- Compute  $h(y)$  for each hash function  $y$ .
- If all the resulting bit positions are 1, say we have seen  $y$  before.
  - We could be wrong.
    - Different inputs could have set each of these bits.
- If at least one of these positions is 0, say we have not seen  $y$  before.
  - We are certainly right.


## Example: Lookup

- Suppose we have the same Bloom filter as before, and we have set the filter to 10100101010.
- Lookup element  $y = 118 = 1110110$  (binary).
- $h_1(y) = 14 \text{ modulo } 11 = 3$ .
- $h_2(y) = 5 \text{ modulo } 11 = 5$ .
- Bit 5 is 1, but bit 3 is 0, so we are sure  $y$  is not in the set.

# Performance of Bloom Filters

- Probability of a false positive depends on the density of 1's in the array and the number of hash functions.
  - $= (\text{fraction of 1's})^{\# \text{ of hash functions}}$ .
- The number of 1's is approximately the number of elements inserted times the number of hash functions.
  - But collisions lower that number slightly.

# Sampling a Stream



# When Sampling Doesn't Work

- Suppose Google would like to examine its stream of **search queries** for the past month to find out what fraction of them were **unique** – asked only once.
- But to save time, we are only going to sample 1/10th of the stream.
- The fraction of unique queries in the sample != the fraction for the stream as a whole.
  - In fact, we can't even adjust the sample's fraction to give the correct answer.



## Example: Unique Search Queries

- The length of the sample is 10% of the length of the whole stream.
- Suppose a query is unique.
  - It has a 10% chance of being in the sample.
- Suppose a query occurs exactly twice in the stream.
  - It has an 18% chance of appearing exactly once in the sample.
- And so on ... The fraction of unique queries in the stream is unpredictably large.

# Sampling by Value

- **My mistake:** I sampled based on the position in the stream, rather than the value of the stream element.
- **The right way:** hash search queries to 10 buckets 0, 1, ..., 9.
- Sample = all search queries that **hash to bucket 0**.
  - All or none of the instances of a query are selected.
  - Therefore the fraction of unique queries in the sample is the same as for the stream as a whole.

# Controlling the Sample Size

- **Problem:** What if the total sample size is limited?
- **Solution:** Hash to a large number of buckets.
- Adjust the set of buckets accepted for the sample, so your sample size stays within bounds.

## Example: Fixed Sample Size

- Suppose we start our search-query sample at 10%, but we want to limit the size.
- Hash to (say) 100 buckets, 0, 1, ..., 99.
  - Take for the sample those elements hashing to buckets 0 through 9.
- If the sample gets too big, get rid of bucket 9.
- Still too big, get rid of 8, and so on.

# Sampling Key-Value Pairs

- This technique generalizes to any form of data that we can see as tuples  $(K, V)$ , where  $K$  is the “key” and  $V$  is a “value.”
- **Distinction:** We want our sample to be based on picking some set of keys only, not pairs.
  - In the search-query example, the data was “all key.”
- Hash keys to some number of buckets.
- Sample consists of all key-value pairs with a key that goes into one of the selected buckets.

## Example: Salary Ranges

- Data = tuples of the form (EmpID, Dept, Salary).
- **Query:** What is the average range of salaries within departments?
- Key = Dept.
- Value = (EmpID, Salary).
- Sample picks some departments, has salaries for all employees of that department, including its min and max salaries.
- Result will be an unbiased estimate of the average salary range.

# Counting Distinct Elements

- **Problem:** a data stream consists of elements chosen from a set of size  $n$ . Maintain a count of the number of distinct elements seen so far.
- **Obvious approach:** maintain the set of elements seen.

# Application Examples

- How many **different words** are found among the Web pages being crawled at a site?
  - Unusually low or high numbers could indicate artificial pages (spam?).
- How many **unique users** visited Facebook this month?
- How many **different pages** link to each of the pages we have crawled.
  - Useful for estimating the PageRank of these pages.
    - Which in turn can tell a crawler which pages are most worth visiting.



# Estimating Counts

- **Real Problem:** what if we do not have space to store the complete set?
  - Or we are trying to count lots of sets at the same time.
- Estimate the count in an unbiased way.
- Accept that the count may be in error, but limit the probability that the error is large.

# Flajolet-Martin Approach



- A probabilistic counting algorithm
- Used to estimate number of **distinct** elements in a large file originally
- Use little memory
- Single pass only
- Based on statistical observation made on bits of hashed values

## Flajolet-Martin Approach (2)

Estimating the cardinality of a multiset  $M$ :

**for**  $i := 0$  **to**  $N - 1$  **do**  $BITMAP[i] := 0$ ;

**for all**  $x$  in  $M$  **do**

**begin**

$index := p(hash(x))$ ;

**if**  $BITMAP[index] = 0$  **then**  $BITMAP[index] := 1$ ;

**end**;

$R :=$  the largest  $index$  in  $BITMAP$  whose value equals to 1

**Estimate**  $:= 2^R$

# Flajolet-Martin Approach (3)



- If the final BITMAP looks like this:  
0000,0000,1100,1111,1111,1111
- The left most 1 at position 15
- Thus, have around  $2^{15}$  distinct elements in the stream