

Recommender systems

Goal

- To know the motivation behind recommender systems (RS)
- To understand how RS works and know the principle behind RS methods/algorithms

Motivation



Problem: How to recommend items to users to make user, content partner, websites happy!

Recommendations



movielens
helping you find the right movies

LinkedIn

clicktorch
Intelligent Recommendation

NETFLIX

facebook

You Tube

amazon

P

Sugestio

last.fm

Spotify

Google news

ebay

From Scarcity to Abundance

- Shelf space scarce commodity for traditional retailers
 - Also: TV networks, movie theatres,...
- Web enables near-zero-cost dissemination of information about products
 - From scarcity to abundance
- More choice necessitates better filters
 - Recommendation engines
 - How Into Thin Air made Touching the Void a bestseller: <http://www.wired.com/wired/archive/12.10/tail.html>

The Long Tail Phenomenon



Types of Recommendations

- Editorial and hand curated
 - List of favorites
 - Lists of “essential” items
- Simple aggregates
 - Top 10, Most Popular, Recent Uploads
- **Tailored to individual users**
 - Amazon, Netflix, Apple Music, Spotify, Facebook Ads...



Formal Model

- **X** = Set of **C**ustomers
- **S** = Set of **I**tems
- **Utility function $u: X \times S \rightarrow R$**
 - **R** = Set of **r**atings
 - **R** totally ordered set
 - E.g., 0 - 5 stars. real number in $[0, 1]$

Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0,2	
Bob		0,5		0,3
Carol	0,2		1	
David				0,4

Knocked Up	🚫 ⭐⭐⭐⭐☆
Evel	🚫 ⭐⭐⭐⭐☆
Dreamgirls	🚫 ⭐⭐⭐⭐☆
The Bridge	🚫 ⭐⭐⭐⭐☆
Children of Men	🚫 ⭐⭐⭐⭐☆
Breach	🚫 ⭐⭐⭐⭐☆
Sweet Land	🚫 ⭐⭐⭐⭐☆
The Good Shepherd	🚫 ⭐⭐⭐⭐☆
Live Free or Die Hard	🚫 ⭐⭐⭐⭐☆
Zodiac	🚫 ⭐⭐⭐⭐☆

Key Problems

1. **Gathering “known” ratings for matrix**

- How to collect the data in the utility matrix

2. **Extrapolate unknown ratings from the known ones**

- Mainly interested in high unknown ratings
- Not interested in knowing what you don't like but what you like

3. **Evaluating extrapolation methods**

- How to measure success/performance of recommendation methods

(1) Gathering Ratings

- **Explicit**
 - Ask people to rate items
 - Doesn't work well in practice – people can't be bothered
 - Crowdsourcing: Pay people to label items
- **Implicit**
 - Learn ratings from user actions
E.g., purchase implies high rating
 - What about low ratings?

(2) Extrapolating Utilities

- **Key problem:**
 - **Sparsity** of utility matrix U
 - Most people have not rated most items
 - **Cold start:**
 - New items - no ratings
 - New users - no history
- **Three approaches to recommender systems:**
 - 1) Content-based
 - 2) Collaborative
 - 3) Latent factor based

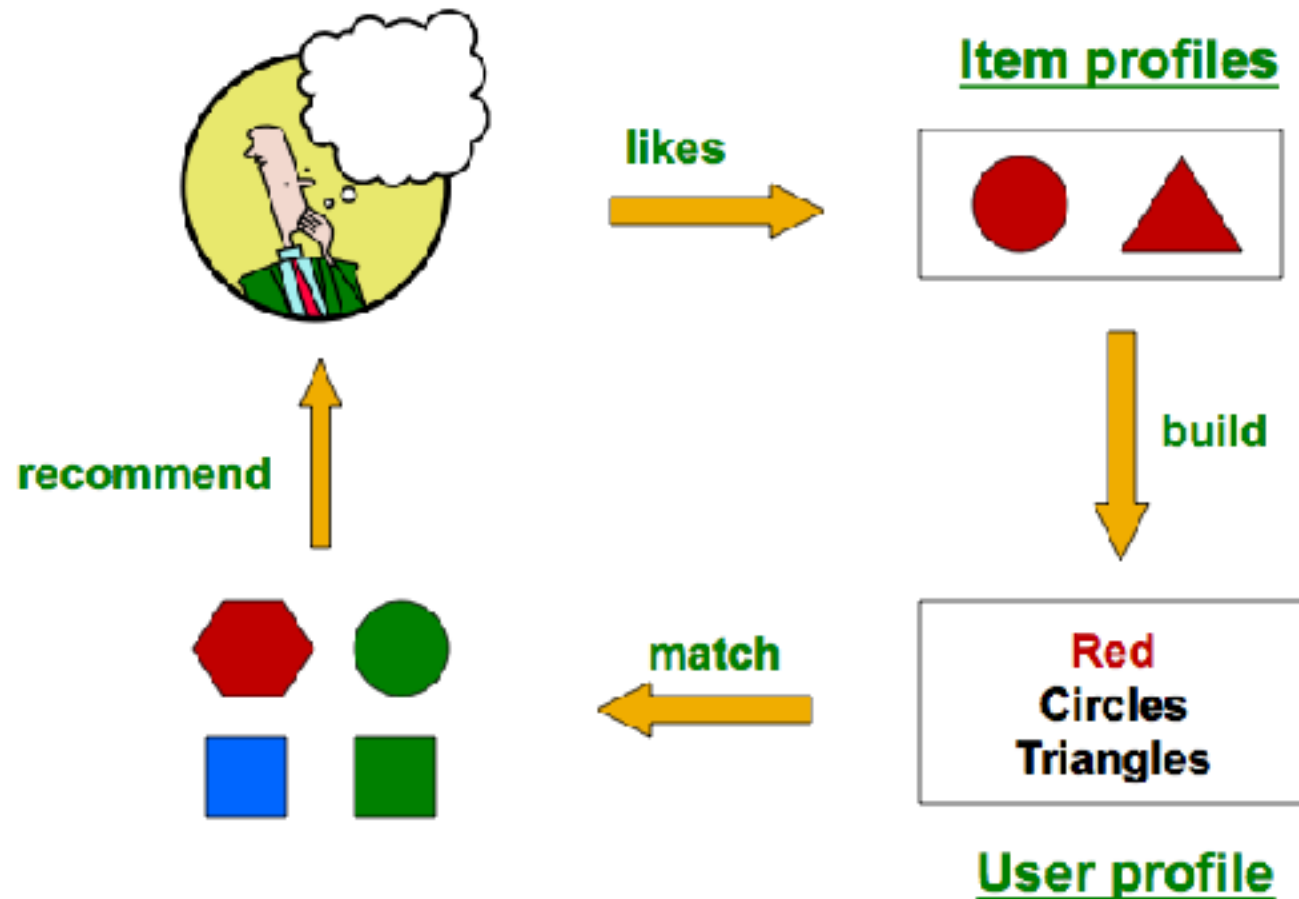
Content-based Recommendations

- **Main idea:**
Recommend items to customer X similar to previous items rated highly by X

Example:

- **Movie recommendations**
 - Recommend movies with same actor(s), director, genre, ...
- **Websites, blogs, news**
 - Recommend other sites with “similar” content

Plan of Action



Item Profiles

- For each item, create an item profile
- Profile: a set (vector) of features
 - **Movies:** author, title, actor, director,...
 - **Text:** Set of “important” words in document
- How to pick important features?
 - Usual heuristic from text mining is TF-IDF (Term frequency * Inverse Doc Frequency)
 - Term ... Feature
 - Document ... Item

Item Features

- Description of items in terms of attributes
 - E.g.: Type, Director, Actors
- Description via keywords
- Possibility to look at content itself, like the text



TF-IDF

f_{ij} = frequency of term (feature) i in doc (item) j

$$TF_{ij} = \frac{f_{ij}}{\max_k f_{kj}}$$

Note: we normalize TF to discount for “longer” documents

n_i = number of docs that mention term i

N = total number of docs

$$IDF_i = \log \frac{N}{n_i}$$

TF-IDF score: $w_{ij} = TF_{ij} \times IDF_i$

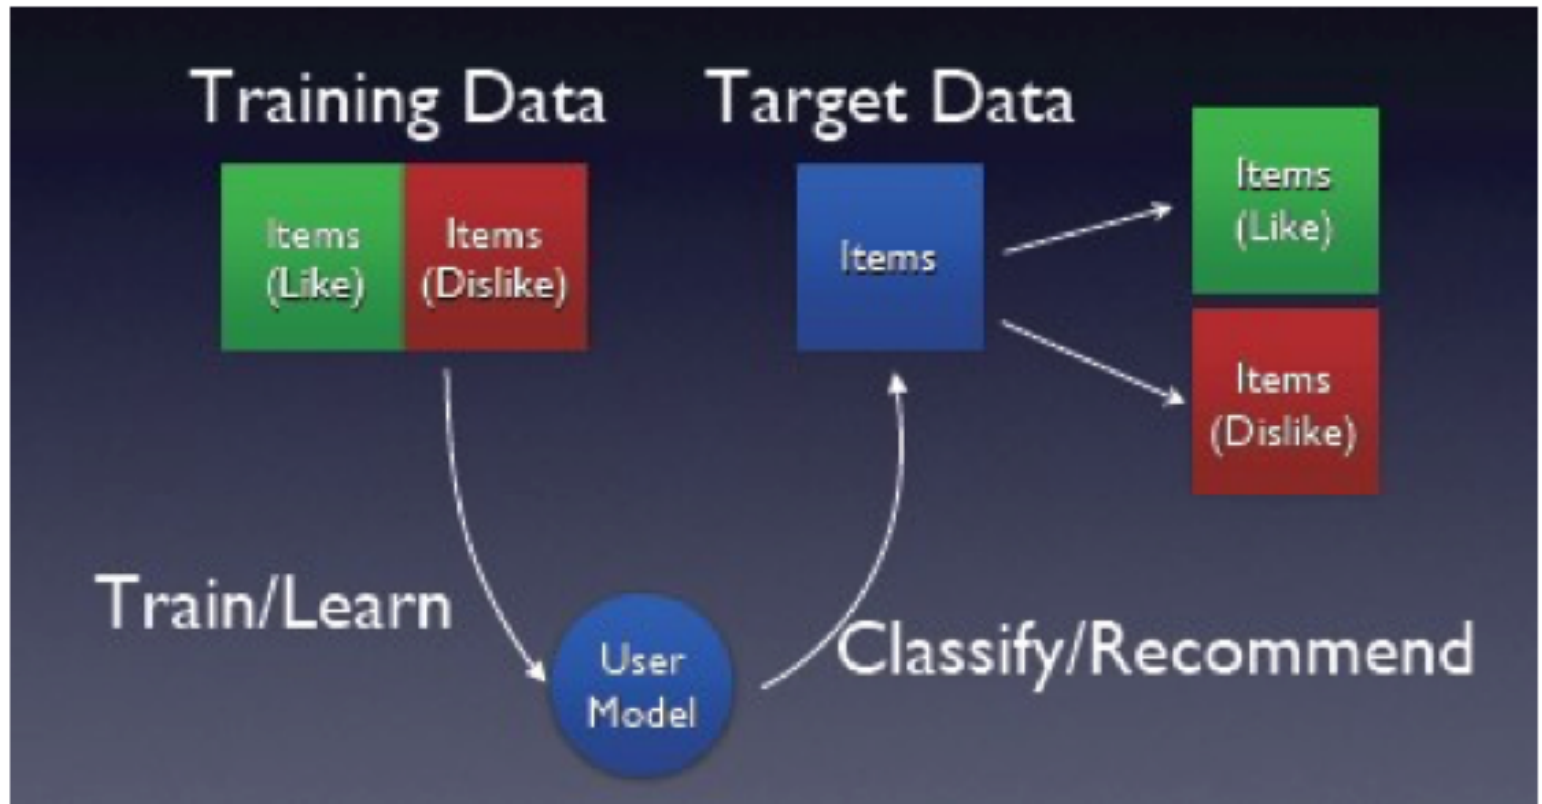
Doc profile = set of words with highest **TF-IDF** scores, together with their scores

User Profiles and Prediction

- User profile possibilities:
 - Weighted average of rated item profiles
 - Variation: weight by difference from average rating for item
 - ...
- Prediction heuristic
 - Given user profile x and item profile i , estimate

$$u(x, i) = \cos(x, i) = \frac{x \cdot i}{||x|| \cdot ||i||}$$

Learning a User Model



Pros: Content-based approach

- **+: No need for data on other users**
 - No cold-start or sparsity problems
- **+: Able to recommend to users with unique tastes**
- **+: Able to recommend new & unpopular items**
 - No first-rater problem
- **+: Able to provide explanations**
 - Can provide explanations of recommended items by listing content-features that caused an item to be recommended

Cons: Content-bases Approach


- **–: Finding the appropriate features is hard**
 - E.g., images, movies, music
- **–: Recommendations for new users**
 - How to build a user profile?
- **–: Overspecialization**
 - Never recommends items outside user's content profile
 - People might have multiple interests
 - **Unable to exploit quality judgments of other users**

Collaborative Filtering

How can we exploit quality judgments of other users to provide relevant recommendation?

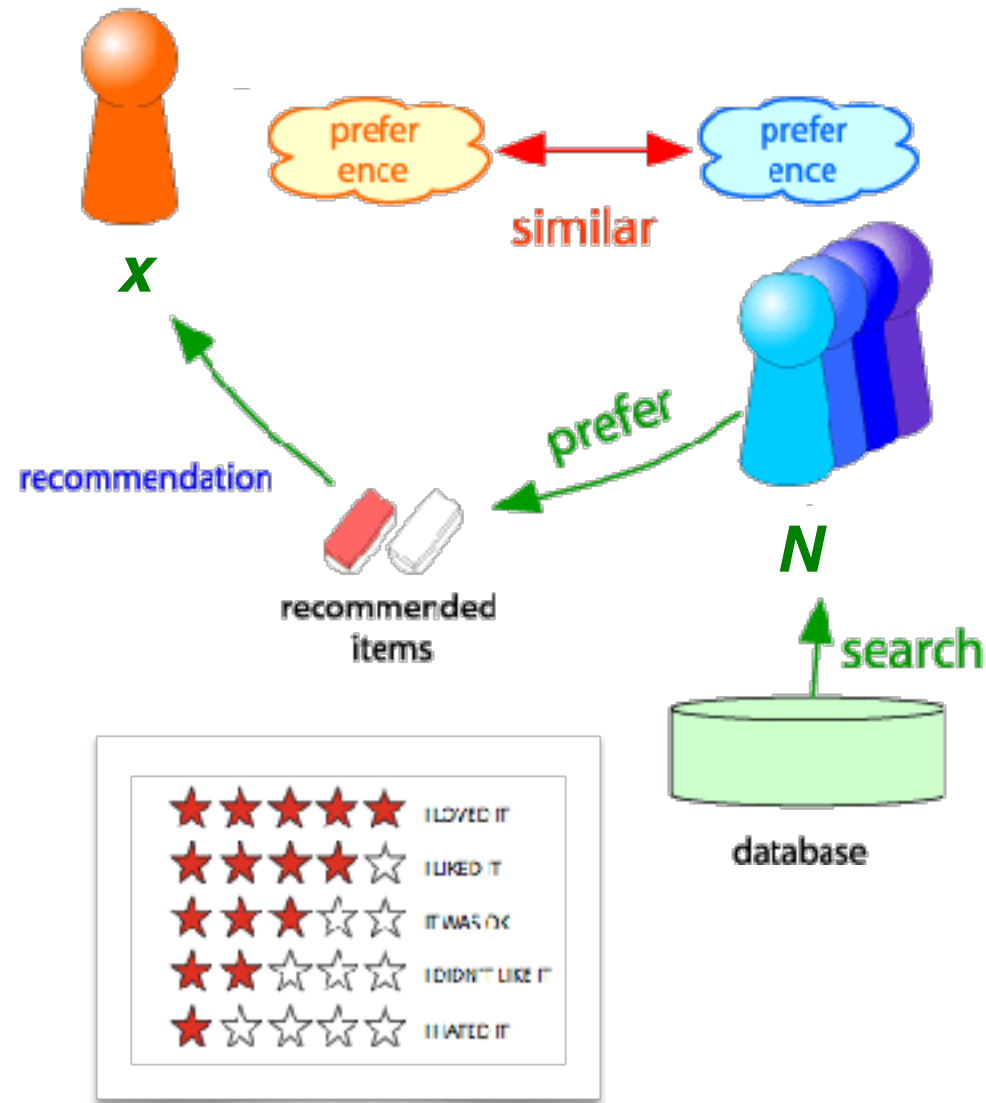
Collaborative Filtering

- Maintain a database of many **users' ratings of a variety of items.**
- For a given user, **find other similar users whose ratings strongly correlate** with the current user.
- **Recommend items rated highly by these similar users,** but not rated by the current user.
- Almost all existing commercial recommenders use this approach (e.g. Amazon).



Collaborative Filtering

- Consider user x
- Find set N of other users whose ratings are “similar” to x ’s ratings
- Estimate x ’s ratings based on ratings of users in N



Find “similar users”

$$r_x = \begin{bmatrix} * & _ & _ & * & *** \end{bmatrix}$$

$$r_y = \begin{bmatrix} * & _ & ** & ** & _ \end{bmatrix}$$

- Let r_x be the vector of user x 's ratings
- Jaccard similarity measure**
 - Problem:** Ignores the value of the rating

r_x, r_y as sets:

$$r_x = \{1, 4, 5\}$$

$$r_y = \{1, 3, 4\}$$

- Cosine similarity measure**

$$\text{sim}(x, y) = \cos(r_x, r_y) = \frac{r_x \cdot r_y}{\|r_x\| \cdot \|r_y\|}$$

r_x, r_y as points:

$$r_x = \{1, 0, 0, 1, 3\}$$

$$r_y = \{1, 0, 2, 2, 0\}$$

- Problem:** Treats missing ratings as “negative”
- Pearson correlation coefficient**
 - S_{xy} = Set of items rated by both users x and y

$$\text{sim}(x, y) = \frac{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)(r_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)^2} \sqrt{\sum_{s \in S_{xy}} (r_{ys} - \bar{r}_y)^2}}$$

$\bar{r}_x, \bar{r}_y \dots$ avg.
rating of x, y

Similarity Metric

$$\text{sim}(x, y) = \frac{\sum_i r_{xi} \cdot r_{yi}}{\sqrt{\sum_i r_{xi}^2} \cdot \sqrt{\sum_i r_{yi}^2}}$$

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- Intuitively we want: $\text{sim}(A, B) > \text{sim}(A, C)$
- **Jaccard similarity:** $1/5 < 2/4$
- **Cosine similarity:** $0.380 > 0.322$
 - Considers missing ratings as “negative”
 - Solution: subtract the (row) mean

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	2/3			5/3	7/3		
B	1/3	1/3	-2/3				
C				-5/3	1/3	4/3	
D		0					0

sim A,B vs. A,C:
0.092 > -0.559

Notice cosine sim. is correlation when data is centered at 0

Rating Predictions

- **From similarity metric to recommendations:**
 - Let r_x : vector of user x 's ratings
 - Let N : set of k users most similar to x who have rated item i
- Prediction for item s of user x :
 - $$r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$$
 - $$r_{xi} = \frac{\sum_{y \in N} s_{xy} \cdot r_{yi}}{\sum_{y \in N} s_{xy}}$$
 - Other options?
- Many other tricks possible...

Shorthand:

$$s_{xy} = \text{sim}(x, y)$$

Item-Item Collaborative Filtering

- So far: **User-user collaborative filtering**
- Another view: **Item-item**
 - For item i , find other similar items
 - Estimate rating for item i based on ratings for similar items
 - Can use same similarity metrics and prediction functions as in user-user model

$$r_{xi} = \frac{\sum_{j \in N(i; x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N(i; x)} s_{ij}}$$

s_{ij} ... similarity of items i and j

r_{xj} ... rating of user u on item j

$N(i; x)$... set items rated by x similar to i

Item-Item CF ($|N|=2$)

users

movies

	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3			5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

- unknown rating- rating between 1 to 5

Item-Item CF ($|N|=2$)

users

movies

	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		?	5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	



- estimate rating of movie 1 by user 5

Item-Item CF ($|N|=2$)

users

	1	2	3	4	5	6	7	8	9	10	11	12	$\text{sim}(1,m)$
1	1		3		?	5			5		4		1.00
2			5	4			4			2	1	3	-0.18
3	2	4		1	2		3		4	3	5		<u>0.41</u>
4		2	4		5			4			2		-0.10
5			4	3	4	2					2	5	-0.31
6	1		3		3			2			4		<u>0.59</u>

Neighbor selection:

Identify movies similar to movie 1, rated by user 5

Here we use Pearson correlation as similarity:

1) Subtract mean rating m_i from each movie i

$$m_1 = (1+3+5+5+4)/5 = 3.6$$

row 1: $[-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]$

2) Compute cosine similarities between rows

Item-Item CF ($|N|=2$)

users

	1	2	3	4	5	6	7	8	9	10	11	12	$\text{sim}(1,m)$
1	1		3		?	5			5		4		1.00
2			5	4			4			2	1	3	-0.18
3	2	4		1	2		3		4	3	5		<u>0.41</u>
4		2	4		5			4			2		-0.10
5			4	3	4	2					2	5	-0.31
6	1		3		3			2			4		<u>0.59</u>

Compute similarity weights:

$$s_{1,3}=0.41, s_{1,6}=0.59$$

Item-Item CF ($|N|=2$)

users

movies

	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		?	5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3		?	2			4	

Predict by taking weighted average:

$$r_{1.5} = (0.41 \cdot 2 + 0.59 \cdot 3) / (0.41 + 0.59) = 2.6$$

$$r_{ix} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{jx}}{\sum s_{ij}}$$

CF: Common Practice

Before:

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

- Define similarity s_{ij} of items i and j
- Select k nearest neighbors $N(i; x)$
 - Items most similar to i , that were rated by x
- Estimate rating r_{xi} as the weighted average:

$$r_{xi} = \underbrace{b_{xi}}_{\text{baseline estimate for } r_{xi}} + \frac{\sum_{j \in N(i;x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i;x)} s_{ij}}$$

baseline estimate for r_{xi}

$$b_{xi} = \mu + b_x + b_i$$

μ = overall mean movie rating

b_x = rating deviation of user x

= (avg. rating of user x) - μ

b_i = rating deviation of movie i

Item-Item vs. User-User

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.8	
Bob		0.5		0.3
Carol	0.9		1	0.8
David			1	0.4

- In practice: item-item often works better than user-user
- Why? Items are simpler, users have multiple tastes

Pros/Cons of Collaborative Filtering

- **+ Works for any kind of item**
 - No feature selection needed
- **- Cold Start:**
 - Need enough users in the system to find a match
- **- Sparsity:**
 - The user/ratings matrix is sparse
 - Hard to find users that have rated the same items

Pros/Cons of Collaborative Filtering (2)

- - **First rater:**
 - Cannot recommend an item that has not been previously rated
 - New items, Esoteric items
- - **Popularity bias:**
 - Cannot recommend items to someone with unique taste
 - Tends to recommend popular items

Cold Start

- **New item problem:**
 - small number of users that rated an item, accurate prediction for this item cannot be generated.
- **New user problem:**
 - small number of items rated by a user, it is unlikely that there could be an overlap of items rated by this user and active users. User– to–user similarity cannot be reliably computed.
- **New community problem:**
 - Without sufficient ratings, it's hard to differentiate value by personalized CF recommendations.
- Clear **reward systems** necessary to convince users to vote or rate items.

Possible solution to cold start

- **As the solution for new user problem:**
 - Displaying non-personalized recommendation until the user has rated enough
 - Asking the user to describe their taste in aggregate
 - Asking the user for demographic information and using ratings of other users with similar demographics as recommendations
- **As the solution for new item problem:**
 - Recommending items through non-CF techniques content analysis or metadata
 - Randomly selecting items with few or no ratings and asking user to rate those items.
- **As the solution for new community problem:**
 - Provide ratings incentives to a small “bootstrap” subset of the community, before inviting the entire community.

Data Sparsity & Ratings scarcity

- The ratings matrix is sparse and only a small fraction of all possible user item entries is known.
- Many CF algorithms designed specifically for data sets with many more users than items (e.g., the MovieLens data set has 65,000 users and 5,000 movies).
- CF may be inappropriate in a domain where there are many more items than users.
- Implicit vs. explicit ratings

Evaluation of Collaborative Filtering

- To determine the quality of the predictions and recommendations
 - **Accuracy**
 - Rating accuracy : error between the predicted ratings and the true ratings. **Mean Absolute Error (MAE)** = average absolute difference between the predicted ratings and the actual rating given by a user
 - **Precision**
 - **Rank accuracy**: half-life utility.
 - **Novelty / Serendipity** (Karypis, 2001)
 - **Coverage** (Sarwar, et. al., 2000)
 - **Learning Rate** (Schein, et. al., 2001)
 - **Confidence** (Herlocker, 2000)
 - **User Satisfaction** (Swearingen & Sinha, 2001; Dahlen, B. J., 1998)
 - Site Performance

Hybrid Methods

- **Implement two or more different recommenders and combine predictions**
 - Perhaps using a linear model
- **Add content-based methods to collaborative filtering**
 - Item profiles for new item problem
 - Demographics to deal with new user problem